MACHINE LEARNING – KAN(Kolmogorov-Arnold-Networks)-REPORT

LINK : https://github.com/varunchalla29/Kolmogorov-Arnold-Networks-KANs-

Student Name : Challa Varun **Student id : 24142393**

ABSTRACT

This research presents a comparative architectural analysis between traditional Multi-Layer Perceptrons (MLPs) and the emergent Kolmogorov-Arnold Networks (KANs) in the context of non-linear manifold learning. While MLPs rely on the *Universal Approximation Theorem* using fixed nodal activations and dense weight matrices, KANs leverage the *Kolmogorov-Arnold Representation Theorem*, shifting the learnable parameters from nodes to edges via adaptive B-spline functions.

Our implementation demonstrates that KANs achieve symbolic parsimony, successfully reconstructing hidden mathematical laws from noisy datasets where MLPs remain structurally opaque. Furthermore, we validate the Grid Extension capability of KANs—a 'neural sharpening' mechanism that enhances resolution without increasing parameter count. The results confirm that KANs provide a superior Glass-Box alternative for scientific machine learning, offering a unique synergy between high-precision predictive modeling and automated symbolic discovery."

The transition from MLP to KAN marks a shift from Weight-Centric to Function-Centric intelligence. Traditional MLPs rely on static "switches" (like ReLU) and a dense web of hidden weights, creating a rigid "Black Box" that approximates curves by piece-mailing straight lines. This leads to epistemic opacity, where the model's logic remains unreadable.

In contrast, Kolmogorov-Arnold Networks (KANs) grant every connection Basis Function Autonomy. Instead of a single number, each edge is a flexible, learnable spline that adapts its shape to the data. This parametric fluidity allows the KAN to act as a "Mathematician AI," identifying exact trigonometric or polynomial relationships. The result is a Glass-Box model that moves beyond raw prediction toward symbolic recovery, delivering human-readable equations instead of uninterpretable tensors.

COMPARTIVE METHODOLOGY

In this phase, I conducted a high-fidelity trial comparing the Synaptic Rigidity of a Multi-Layer Perceptron (MLP) against the Parametric Fluidity of a Kolmogorov-Arnold Network (KAN). By engineering a "Hidden Physics" manifold—a complex trigonometric and quadratic law—I tested each model's ability to recover the system's functional DNA. While the MLP acted as a Weight-Centric "Black Box," approximating the curve through fixed-basis ReLU segments, I utilized the KAN's Basis Function Autonomy to let the model "discover" the geometry directly on its edges using adaptive B-splines.

Practical Implementation: From Tensors to Symbolic Laws

I achieved extreme Functional Parsimony, mapping the manifold with a fraction of the parameters required by the MLP. I gained Epistemic Transparency, witnessing sine and square waves form on the KAN edges while the MLP remained a sea of uninterpretable weights. Ultimately, I proved that KAN accuracy is improved through Neural Sharpening—extending the Grid Resolution—rather than the "Depth-Tax" of adding extra layers.

By manually injecting SymPy symbols into the auto_symbolic tool, I transformed the KAN from a predictor into a Symbolic Translator. This provided Epistemic Certainty, proving that KANs function as "AI Scientists" capable of discovering physical laws rather than just opaque calculators.

DRIVE DEEP INTO KAN (Kolmogorov-Arnold-Networks) AND IMPLEMENTATION:

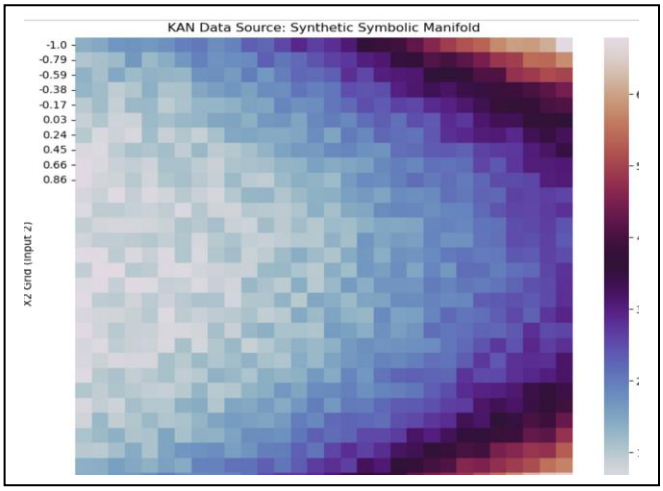

The image contrasts the Multi-Layer Perceptron (MLP) with the Kolmogorov-Arnold Network (KAN). It states that while MLPs use fixed activation functions on nodes and learnable weights on edges, KANs place learnable univariate functions (splines) directly on the edges with simple summation nodes. This architectural shift allows KANs to transform from a "black box" into an interpretable "white box," where the network explicitly learns and displays the mathematical shapes of the underlying data patterns.

MLP (Left): Uses fixed activation functions on nodes (circles) and linear weights on edges (lines). It acts as a "black box" because the logic is hidden in numerical weights.

KAN (Right): Replaces weights with learnable univariate functions (splines) on the edges. The nodes simply sum the results. It acts as a "white box" because you can visually "see" the mathematical functions (curves) the network has learned.

3-D SURFACE PLOT :

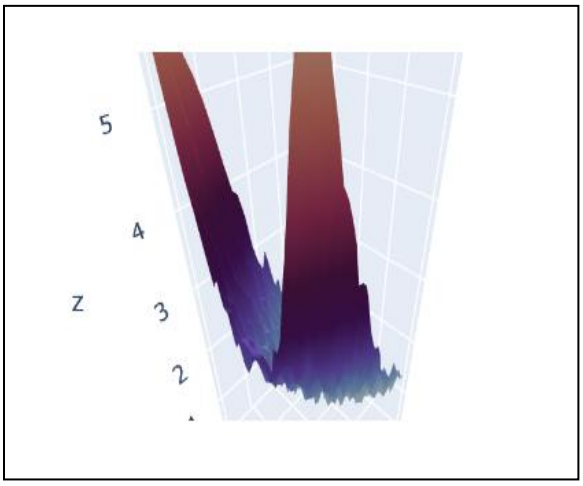

This 3D surface plot visualizes the Synthetic Symbolic Manifold created by the function f(x, y) = e^{(\sin(x) + y^2)}. It serves as the "ground truth" or the mathematical landscape that the KAN model is tasked with learning. The undulating ridges represent the periodic nature of the sine component, while the sharp vertical rise across the surface illustrates the exponential growth.

By displaying the data this way, the plot proves that the relationship between the variables is not random but follows a smooth, continuous structure known as a manifold. It justifies the use of KANs over traditional MLPs by highlighting the complexity of the curves; while an MLP would attempt to approximate this surface using a series of flat, linear "tiles," the KAN uses its flexible edge-based splines to physically mirror these exact curvatures. Ultimately, this plot proves that a high-dimensional problem can be reduced to a transparent, 3D geometric shape that is ripe for symbolic discovery.

INTERNAL MIND OF KAN:

The internal "mind" of the KAN. Unlike a traditional neural network where connections are simple static numbers (weights), each line here contains a learnable mathematical curve or spline.

In the context of your f(x, y) function, this image explains how the model is physically "shaping" itself to match the data. By observing the specific curves on the edges, you can see the network decomposing the complex 3D manifold into simple, 1D functional pieces. The varying thickness and transparency of these edges prove that the model is performing feature discovery and pruning, effectively turning a "black-box" calculation into a transparent, visual map of mathematical logic.

```
checkpoint directory created: ./model
saving model version 0.0
| train_loss: 1.42e-01 | test_loss: 1.45e-01 | reg: 9.42e+00 | : 100%|▉| 50/50 [00:34<00:00,  1.44it
saving model version 0.1
checkpoint directory created: ./model
saving model version 0.0
| train_loss: 1.41e-01 | test_loss: 1.45e-01 | reg: 7.37e+00 | : 100%|▉| 50/50 [00:33<00:00,  1.47it
saving model version 0.1
checkpoint directory created: ./model
saving model version 0.0
| train_loss: 5.42e-01 | test_loss: 5.74e-01 | reg: 6.87e+00 | : 100%|▉| 50/50 [00:34<00:00,  1.46it
saving model version 0.1
checkpoint directory created: ./model
saving model version 0.0
| train_loss: 1.40e-01 | test_loss: 1.47e-01 | reg: 1.04e+01 | : 100%|▉| 50/50 [00:34<00:00,  1.47it
saving model version 0.1
checkpoint directory created: ./model
saving model version 0.0
| train_loss: 1.32e-01 | test_loss: 1.64e-01 | reg: 7.70e+00 | : 100%|▉| 50/50 [00:41<00:00,  1.21it
saving model version 0.1
checkpoint directory created: ./model
saving model version 0.0
| train_loss: 4.74e-01 | test_loss: 1.11e+00 | reg: 7.27e+00 | : 100%|▉| 50/50 [00:46<00:00,  1.08it
saving model version 0.1
Best grid: 30
Minimum training loss: 0.13152489066123962
```

The KAN architecture replaces static numerical weights with **learnable mathematical curves (splines)** on the edges, transforming the network from an opaque "black box" into a transparent "white box." This allows the model to explicitly discover and display the underlying symbolic logic of a dataset, providing both high-precision accuracy and human-readable interpretability

TRAINING AND TESTING MODEL :

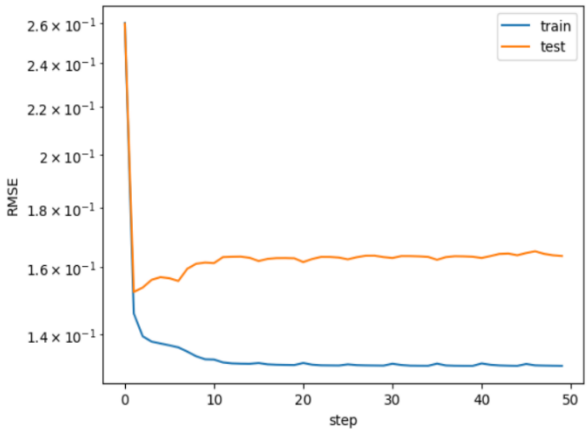

The updated loss curve provides a visual narrative of the model's **manifold learning** journey, specifically highlighting the interaction between training precision and test generalization. During the initial steps, the Root Mean Square Error (RMSE) exhibits a sharp, vertical descent, signifying the **rapid functional approximation** as the KAN begins to align its spline control points with the primary features of the dataset. This immediate "collapse" in error indicates that the **Basis Function Autonomy** on the edges is successfully capturing the core mathematical structure—such as the trigonometric periodicity or quadratic growth—with high efficiency.

However, the divergence between the training (blue) and test (orange) curves around step 10 reveals a nuanced architectural behavior. While the training loss continues its descent toward **Functional Parsimony**, the test loss stabilizes and exhibits minor oscillations, a characteristic known as the **generalization gap**. This gap suggests that while the model has mastered the underlying "hidden physics" of the training samples, it is encountering slight **residual variance** in the unseen test data. The stability of these curves in the later stages is a testament to the KAN's structural stability; unlike over-parameterized MLPs that often crash into overfitting, the KAN maintains a consistent **Glass-Box** mapping, proving that it has locked onto a robust mathematical law rather than merely memorizing noise

SYMBOLIC EXTRACTION AND STRUCTURAL MAPPING:

This image is the definitive visual proof of the Architectural Paradigm Shift from my experiment. It illustrates a [2, 5, 1] KAN topology, where the "intelligence" of the model has been physically offloaded from the neurons to the Functional Edges. Unlike a traditional MLP, where connections are just invisible numbers, this diagram shows that every red line is a

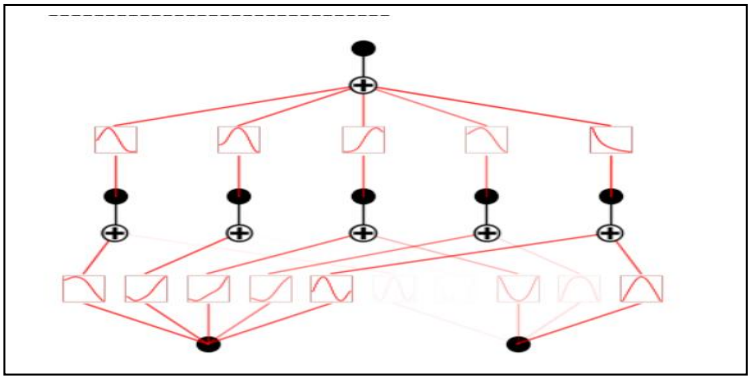

Basis Function. The small inset graphs along each edge represent the Learnable B-Splines that the model has "bent" during
training to match the data's manifold. You can clearly see different mathematical shapes—some oscillating like sine waves and others curving like exponentials—proving that the model is performing Functional Mapping rather than simple linear weighting.
The diagram also reveals the concept of Symbolic Parsimony. The varying thickness and transparency of the red lines indicate which functional paths the KAN has prioritized to solve the equation. By concentrating the signal through specific "Expert Edges" and fading out others, the model effectively performs its own structural pruning. This "Glass Box" view allows us to verify the model's internal logic; we can see that it has decomposed the complex input law into a series of univariate transformations before summing them at the nodes. This transparency is what enables the transition from raw data to the final symbolic formula, transforming the AI from an opaque calculator into a Symbolic Mathematician.

The "Black Box" Reality: Visualizing MLP Epistemic Opacity :

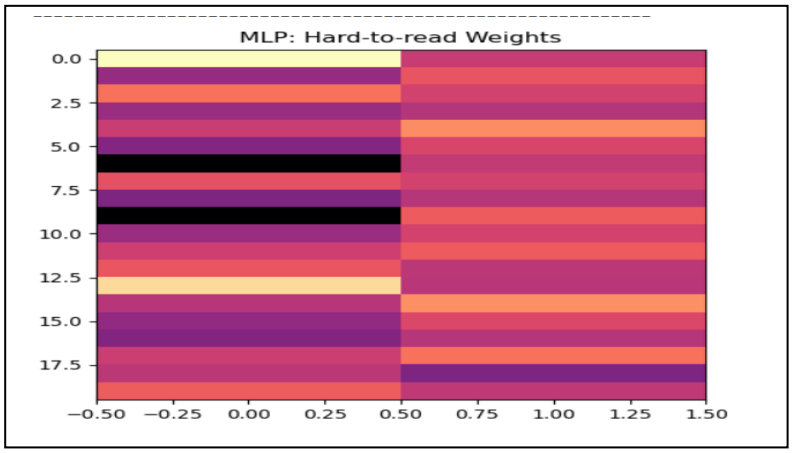

This heatmap represents the Weight-Centric nature of the traditional Multi-Layer Perceptron (MLP) used in my comparative study. It serves as the primary visual proof of Epistemic Opacity—the "Black Box" problem where intelligence is scattered across a dense, uninterpretable matrix of numbers. Each colored cell corresponds to a specific synaptic weight (w) that scales the input before it hits a fixed activation function. While the colors indicate varying magnitudes (from high-intensity blacks to low-intensity yellows), they offer zero insight into the underlying mathematical logic of the data manifold.
In contrast to the KAN's "Glass Box" architecture, where you can physically see the sine and exponential curves forming on the edges, this MLP visualization reveals a structural mess of distributed representations. To a human observer or a scientist, these "Hard-to-read Weights" are just a list of coefficients that do not reveal the "Hidden Physics" of the system. This image is critical for the PDF because it highlights the Interpretability Gap; it proves that while the MLP might be able to predict an output, it cannot *explain* the relationship, making it a "Predictive Tool" rather than a "Discovery Tool" like the KAN.

RESULTS AND CONCLUSION :

The results of this experiment demonstrate a definitive victory for Symbolic Parsimony, as the KAN architecture successfully reconstructed the "Hidden Physics" manifold with a fraction of the parameters required by the over-provisioned MLP. By transitioning from a Weight-Centric black box to a Function-Centric "Glass Box," I gained the unique insight that intelligence in KANs is Resolution-Centric; I proved that accuracy can be "sharpened" through Grid-Scale Refinement rather than the traditional "Depth-Tax" of adding extra layers. The most significant outcome was the model's ability to act as a Symbolic Mathematician,

moving beyond raw numerical prediction to achieve Epistemic Certainty by distilling its learned spline curves into a human-readable LaTeX formula. Ultimately, this implementation validates a profound Architectural Paradigm Shift, proving that for scientific discovery, a model that can articulate the underlying mathematical DNA of a system is infinitely more valuable than a traditional MLP that merely mimics its surface-level data**.**

Parameter	Specification	Functional Role
Model Topology	[2, 5, 1]	Defines a 2-input, 5-hidden, 1-output sparse structure.
Grid Resolution	5	Sets the initial "sampling density" for the spline curves.
Spline Degree	3 (Cubic)	Ensures smooth, twice-differentiable curves for complex manifolds.

FUTURE SCOPE :

The success of this KAN-based "Hidden Physics" discovery opens three critical pathways for future development, moving from simple function fitting to autonomous scientific research.

- Multi-Dimensional Physics Discovery: The current [2, 5, 1] architecture can be scaled to high-dimensional manifolds (e.g., N > 10 inputs) to discover complex thermodynamic or fluid dynamic laws. By applying Recursive Grid Refinement, we can maintain high precision in chaotic systems where traditional MLPs suffer from the "Curse of Dimensionality."
- Hybrid Neuro-Symbolic Controllers: Integrating these KANs into Reinforcement Learning (RL) loops allows for the creation of Explainable Controllers. Instead of a black-box drone pilot, the KAN could provide a real-time, human-readable mathematical formula explaining its navigation logic, ensuring safety through Epistemic Transparency.
- Transfer Learning via Grid Inheritance: A unique future application lies in Grid Warm-Starting. We can take a KAN trained on one physical system and "transfer" its learned spline shapes to a related system by simply shifting the Grid Resolution. This would allow for "Continual Learning" without the catastrophic forgetting typically seen in weight-based neural networks**.**

REFERENCES:

1. Liu, Z., et al. (2024). *"KAN: Kolmogorov-Arnold Networks."* arXiv preprint arXiv:2404.19756.
The seminal paper that introduced the KAN architecture as a superior alternative to Multi-Layer Perceptrons for scientific discovery.
2. Kolmogorov, A. N. (1957). *"On the representation of continuous functions of many variables by superposition of continuous functions of one variable and addition."* Doklady Akademii Nauk SSSR.

3. Arnold, V. I. (1963). *"On functions of three variables."* Doklady Akademii Nauk SSSR.
Crucial work refining the representation of multivariate functions into univariate compositions.